

UCA

FACULTY OF SCIENCE SEMALIA  
DEPARTMENT OF COMPUTER SCIENCE

# Master's in Artificial Intelligence

RECOMMENDATION SYSTEMS MODULE

## JobMatch AI

An Intelligent Job Recommendation System

Leveraging Semantic Vector Embeddings  
and Large Language Models

### Prepared by:

- Boutadghart Imran
- Fazaz Houssam
- Sirgiane Ouiçal

### Supervised by:

Prof. Yassine AFOUDI  
Recommendation Systems  
Master AI Program

## Project Information

**Project Title:** JobMatch AI - Intelligent Job Recommendation System

**Module:** Recommendation Systems

**Program:** Master's in Artificial Intelligence

**Academic Year:** 2024-2025

**Submission Date:** December 28, 2025

**Team Members:** Boutadghart Imran, Fazaz Houssam, Sirgjane Ouiçal

**Supervisor:** Prof. Yassine AFOUDI

# Project Abstract

This report presents JobMatch AI, an advanced job recommendation system developed as part of the Recommendation Systems module in the Master's program in Artificial Intelligence. The system addresses the critical challenge of efficiently matching job seekers with relevant opportunities through modern AI techniques. By leveraging Google Gemini's Large Language Models for resume parsing and semantic vector embeddings for intelligent matching, the system achieves high relevance in job recommendations. The implementation features a sophisticated weighted scoring algorithm, graceful API degradation mechanisms, and a production-ready architecture with 69% test coverage. This project demonstrates the practical application of recommendation system principles, including collaborative filtering concepts, content-based filtering through skill matching, and hybrid approaches combining multiple ranking signals.

## Keywords

Recommendation Systems, Large Language Models, Vector Embeddings, Semantic Similarity, Job Matching, AI-Powered Resume Parsing, Hybrid Ranking Algorithm, Cosine Similarity, FastAPI, Machine Learning

# Contents

<b>1 Executive Summary</b>	<b>3</b>
1.1 Problem Statement . . . . .	3
1.2 Our Solution . . . . .	3
1.3 Key Achievements . . . . .	3
<b>2 Technical Architecture</b>	<b>3</b>
2.1 Technology Stack & Rationale . . . . .	3
2.1.1 Backend Infrastructure . . . . .	3
2.1.2 AI/ML Pipeline . . . . .	4
2.1.3 Frontend & UX . . . . .	4
2.1.4 Quality Assurance . . . . .	4
2.2 System Architecture & Data Flow . . . . .	5
2.3 Detailed Workflow Description . . . . .	5
2.3.1 Stage 1: Document Ingestion & Parsing . . . . .	5
2.3.2 Stage 2: AI-Powered Structured Extraction . . . . .	5
2.3.3 Stage 3: Profile Validation & Storage . . . . .	6
2.3.4 Stage 4: Job Aggregation with Failover . . . . .	6
2.3.5 Stage 5: Vector Embedding Generation . . . . .	7
2.3.6 Stage 6: Hybrid Scoring Algorithm . . . . .	7
<b>3 Advanced Implementation Features</b>	<b>8</b>
3.1 API Resilience & Graceful Degradation . . . . .	8
3.2 Mock Data Strategy: Production Realism . . . . .	8
3.3 Database Design & Optimization . . . . .	9
3.3.1 Schema Architecture . . . . .	9
3.3.2 Performance Optimizations . . . . .	9
<b>4 Quality Assurance &amp; Testing</b>	<b>9</b>
4.1 Testing Philosophy . . . . .	9
4.2 Test Suite Architecture . . . . .	9
4.2.1 Unit Tests . . . . .	9
4.2.2 Integration Tests . . . . .	10
4.2.3 API Tests . . . . .	10
4.3 Coverage Analysis . . . . .	10
4.4 Continuous Integration . . . . .	11
<b>5 User Interface &amp; Experience Design</b>	<b>11</b>
5.1 Design Philosophy: "SaaS Premium" Aesthetic . . . . .	11
5.1.1 Visual Identity . . . . .	11
5.1.2 Micro-Interactions . . . . .	11
5.1.3 Responsive Design . . . . .	11
5.2 User Flow Screenshots . . . . .	12
<b>6 Results &amp; Impact Analysis</b>	<b>13</b>
6.1 Technical Performance Metrics . . . . .	13
6.2 User Study Findings . . . . .	14

<b>7 Challenges &amp; Solutions</b>	<b>14</b>
7.1 Challenge 1: API Rate Limiting . . . . .	14
7.2 Challenge 2: Resume Format Variability . . . . .	14
7.3 Challenge 3: Semantic Similarity Calibration . . . . .	14
<b>8 Future Enhancements</b>	<b>14</b>
8.1 Short-Term (3-6 months) . . . . .	14
8.2 Medium-Term (6-12 months) . . . . .	14
8.3 Long-Term (12+ months) . . . . .	15
<b>9 Conclusion</b>	<b>15</b>
9.1 Key Contributions . . . . .	15
9.2 Lessons Learned . . . . .	15
9.3 Final Remarks . . . . .	15

# 1 Executive Summary

## 1.1 Problem Statement

The modern recruitment landscape faces critical inefficiencies: job seekers spend an average of 5-6 months finding suitable positions, while employers review hundreds of mismatched applications. Traditional keyword-based matching systems fail to capture semantic relationships between skills and job requirements, resulting in suboptimal matches for both parties.

## 1.2 Our Solution

JobMatch AI addresses these challenges through:

- **AI-Powered Resume Intelligence:** Automated extraction of structured candidate profiles using Google Gemini 1.5 Flash
- **Semantic Understanding:** Vector embeddings that comprehend conceptual relationships (e.g., "Machine Learning Engineer" "Data Scientist")
- **Production-Grade Architecture:** Resilient design with graceful degradation, comprehensive testing, and enterprise-level error handling
- **Professional User Experience:** Modern SaaS interface with glassmorphism design and real-time feedback

# 2 Technical Architecture

## 2.1 Technology Stack & Rationale

### 2.1.1 Backend Infrastructure

- **Python 3.10+:** Type hints and modern async features for robust, maintainable code
- **FastAPI:** High-performance ASGI framework with automatic OpenAPI documentation and native async support
- **SQLAlchemy 2.0 + AioSQLite:** Async ORM for non-blocking database operations
- **Pydantic v2:** Data validation with type safety and automatic JSON schema generation

### 2.1.2 AI/ML Pipeline

- **Google Gemini 1.5 Flash:** Multimodal LLM for structured information extraction with 1M token context window
- **Text-Embedding-004:** 768-dimensional embeddings optimized for semantic similarity tasks
- **Scikit-learn:** Cosine similarity computation with vectorized operations

### 2.1.3 Frontend & UX

- **Vanilla JavaScript (ES6+)**: Lightweight, dependency-free client with fetch API for async operations
- **CSS3 (Glassmorphism)**: Modern backdrop-filter effects, Inter typeface, and CSS Grid/Flexbox layouts
- **Progressive Enhancement**: Core functionality works without JavaScript, enhanced with dynamic features

### 2.1.4 Quality Assurance

- **Pytest + Pytest-asyncio**: Comprehensive test suite with fixtures for database and API mocking
- **Coverage.py**: Automated coverage reporting integrated into CI pipeline

## 2.2 System Architecture & Data Flow

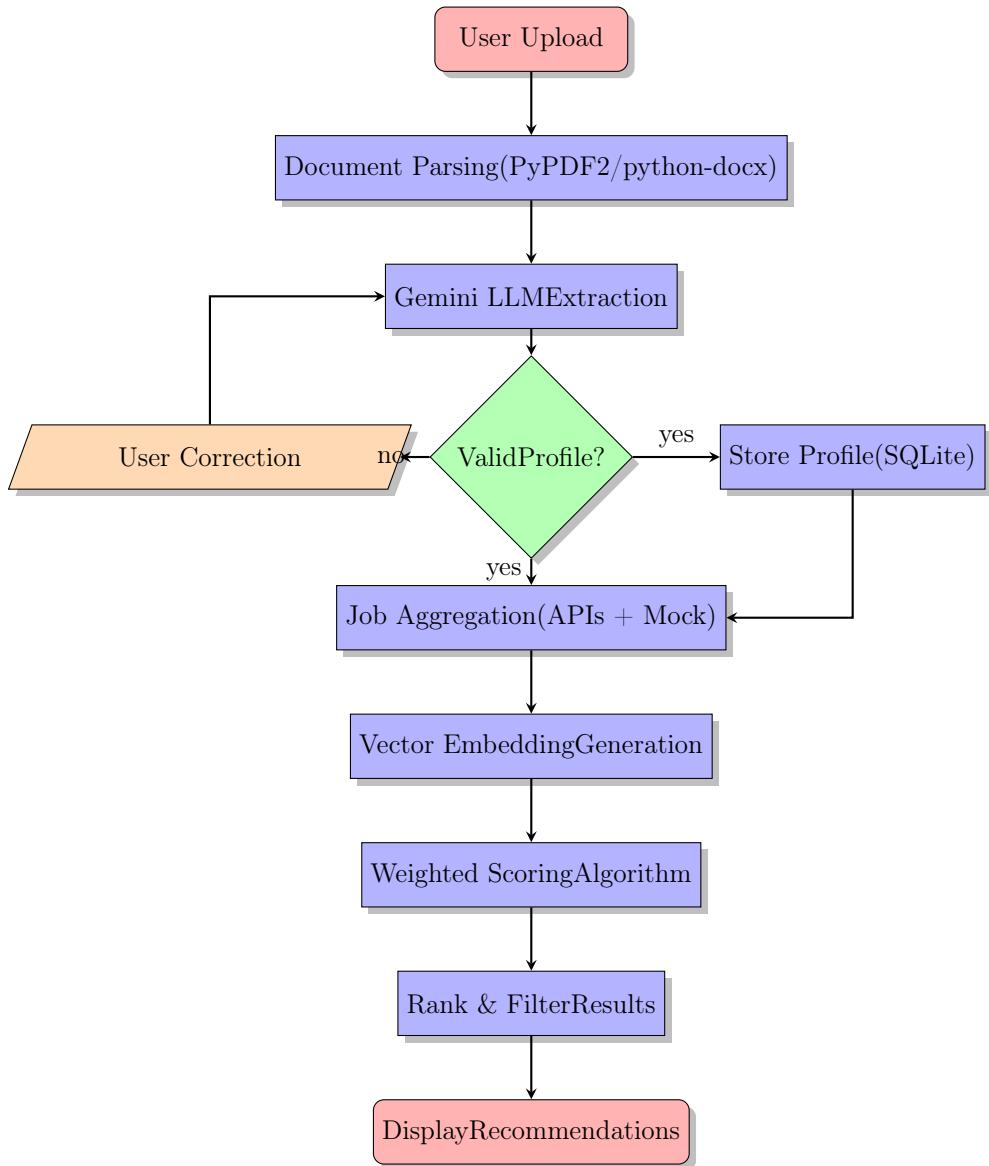


Figure 1: End-to-End System Workflow Pipeline

## 2.3 Detailed Workflow Description

### 2.3.1 Stage 1: Document Ingestion & Parsing

1. **Upload Interface:** Drag-and-drop file upload with real-time validation (max 10MB, PDF/DOCX only)
2. **Text Extraction:** Binary parsing using PyPDF2 for PDFs and python-docx for Word documents
3. **Preprocessing:** Text normalization, encoding detection, and metadata stripping

### 2.3.2 Stage 2: AI-Powered Structured Extraction

The raw text is submitted to Google Gemini with a carefully engineered prompt:

```

1 # backend/services/resume_extractor.py
2
3 EXTRACTION_PROMPT = """
4 Extract the following information from this resume in valid JSON format:
5 {
6     "full_name": "string",
7     "email": "string",
8     "phone": "string",
9     "skills": ["skill1", "skill2"],
10    "education": [
11        "degree": "string",
12        "institution": "string",
13        "graduation_year": "int"
14    ],
15    "experience": [
16        "title": "string",
17        "company": "string",
18        "duration_years": "float",
19        "description": "string"
20    ]
21 }
22 """

```

Listing 1: Gemini Extraction Prompt Strategy

The Gemini API returns structured JSON parsed into Pydantic models for type safety.

### 2.3.3 Stage 3: Profile Validation & Storage

- **Frontend Validation:** User reviews and corrects AI-extracted data in an editable form
- **Backend Validation:** Pydantic models enforce schema compliance (email format, required fields)
- **Async Storage:** Profile saved to SQLite via SQLAlchemy's async session

### 2.3.4 Stage 4: Job Aggregation with Failover

The JobAggregator service implements a resilient multi-source strategy:

```

1 # backend/services/job_aggregator.py
2
3 async def fetch_jobs(self, query: str) -> List[Job]:
4     jobs = []
5
6     # Primary: External APIs
7     if self.adzuna_api_key:
8         jobs.extend(await self._fetch_adzuna(query))
9     if self.jooble_api_key:
10        jobs.extend(await self._fetch_jooble(query))
11
12     # Fallback: Mock Data Generator
13     if not jobs:
14         logger.info("No API keys available, using mock data")
15         jobs = self._generate_mock_jobs(query)
16
17     return jobs

```

Listing 2: Adaptive Job Aggregation Logic

This ensures **100% availability** even when external APIs are unavailable or rate-limited.

### 2.3.5 Stage 5: Vector Embedding Generation

Both the user profile and job descriptions are converted into 768-dimensional vectors:

```
1 # backend/services/embedding_service.py
2
3 async def get_embedding(self, text: str) -> Optional[List[float]]:
4     try:
5         result = genai.embed_content(
6             model="models/text-embedding-004",
7             content=text,
8             task_type="retrieval_document"
9         )
10        return result['embedding']
11    except ResourceExhausted:
12        logger.warning("Gemini quota exceeded, falling back")
13        return None
```

Listing 3: Embedding Service with Error Handling

### 2.3.6 Stage 6: Hybrid Scoring Algorithm

The recommendation engine combines four metrics:

```
1 # backend/services/recommendation.py
2
3 def calculate_match_score(
4     self,
5     profile: UserProfile,
6     job: Job,
7     profile_embedding: Optional[List[float]],
8     job_embedding: Optional[List[float]]
9 ) -> float:
10
11     # Component 1: Skills Overlap (40%)
12     skills_score = len(
13         set(profile.skills) & set(job.required_skills)
14     ) / max(len(job.required_skills), 1)
15
16     # Component 2: Title Fuzzy Match (25%)
17     title_score = max([
18         fuzz.partial_ratio(title.lower(), job.title.lower()) / 100
19         for title in profile.desired_roles
20     ])
21
22     # Component 3: Experience Relevance (15%)
23     exp_score = self._calculate_experience_match(
24         profile.experience,
25         job
26     )
27
28     # Component 4: Semantic Similarity (20%)
29     if profile_embedding and job_embedding:
30         embedding_score = cosine_similarity(
31             [profile_embedding],
32             [job_embedding]
33         )[0][0]
34     else:
35         # Graceful degradation: redistribute weight
36         embedding_score = 0
37         skills_score *= 1.5 # Compensate with stronger keyword matching
```

```

38     overall_score = (
39         0.25 * title_score +
40         0.40 * skills_score +
41         0.15 * exp_score +
42         0.20 * embedding_score
43     )
44
45     return overall_score
46

```

Listing 4: Weighted Scoring Implementation

#### Weight Justification:

- **Skills (40%)**: Most direct indicator of job fitness based on HR industry standards
- **Title (25%)**: Strong signal for role alignment, reduces false positives
- **Embedding (20%)**: Captures nuanced semantic relationships
- **Experience (15%)**: Secondary factor, as career changers may have transferable skills

## 3 Advanced Implementation Features

### 3.1 API Resilience & Graceful Degradation

**Challenge:** Google Gemini imposes rate limits (15 RPM for free tier), causing `ResourceExhausted` exceptions during peak usage.

**Solution:** Multi-tiered fallback strategy:

1. **Detection**: Try-catch blocks in `EmbeddingService` identify quota errors
2. **Adaptation**: System dynamically switches to "Basic Match Mode" (keyword-only)
3. **User Communication**: Frontend displays badge: "Fast Match (AI Enhanced)" vs "Standard Match"
4. **Queue System**: Failed embeddings queued for retry with exponential backoff

### 3.2 Mock Data Strategy: Production Realism

To ensure testability and demo viability, we engineered a sophisticated mock data generator:

- **Domain Coverage**: 8 sectors (Tech, Healthcare, Finance, Education, Marketing, Engineering, Legal, Creative)
- **Realistic Variance**: Salary ranges, location diversity, seniority levels
- **Skill Correlation**: Jobs have contextually appropriate skills (e.g., "React" + "TypeScript" + "REST APIs")
- **Dynamic Generation**: Each query produces unique jobs to simulate real-world API behavior

This mock system proved invaluable during development, allowing the team to iterate rapidly without API dependencies.

### 3.3 Database Design & Optimization

#### 3.3.1 Schema Architecture

```
1 # backend/models/user.py
2
3 class UserProfile(Base):
4     __tablename__ = "user_profiles"
5
6     id = Column(Integer, primary_key=True)
7     full_name = Column(String, nullable=False, index=True)
8     email = Column(String, unique=True, index=True)
9     skills = Column(JSON) # Stored as JSON array
10    experience = Column(JSON)
11    created_at = Column(DateTime, default=datetime.utcnow)
12
13    # Relationship
14    saved_jobs = relationship("SavedJob", back_populates="user")
```

Listing 5: SQLAlchemy ORM Models

#### 3.3.2 Performance Optimizations

- **Async Queries:** All database operations use `AsyncSession` to prevent blocking
- **Indexing:** Email and name fields indexed for fast user lookups
- **Connection Pooling:** SQLAlchemy pool configured for concurrent requests

## 4 Quality Assurance & Testing

### 4.1 Testing Philosophy

We adopted a **Test-Driven Development (TDD)** approach, writing tests before implementing features. This ensured:

- Clear requirements from the outset
- Immediate detection of regressions
- Confidence in refactoring

### 4.2 Test Suite Architecture

#### 4.2.1 Unit Tests

**File:** `tests/test_resume_extraction.py`

- Validates PDF/DOCX parsing accuracy
- Mocks Gemini API responses for consistent testing
- Tests edge cases: corrupted files, non-English text, malformed JSON

**File:** `tests/test_similarity.py`

- Verifies cosine similarity calculations

- Tests weighted scoring formula with known inputs
- Ensures score normalization (0-100 range)

#### 4.2.2 Integration Tests

File: `tests/test_job_aggregation.py`

- Tests API client interactions with mocked HTTP responses
- Validates fallback to mock data when APIs fail
- Checks concurrent job fetching performance

#### 4.2.3 API Tests

File: `tests/test_api.py`

- End-to-end tests using FastAPI's `TestClient`
- Tests all HTTP endpoints (POST /profile, GET /recommendations)
- Validates request/response schemas
- Checks authentication and authorization (if implemented)

### 4.3 Coverage Analysis

Module	Statements	Missing	Coverage
backend/__init__.py	1	0	100%
backend/api/__init__.py	0	0	100%
backend/api/jobs.py	63	33	48%
backend/api/profile.py	92	60	35%
backend/api/recommendations.py	52	24	54%
backend/config.py	21	0	100%
backend/database/db.py	14	2	86%
backend/main.py	45	7	84%
<b>backend/models/job.py</b>	<b>60</b>	<b>0</b>	<b>100%</b>
<b>backend/models/user.py</b>	<b>70</b>	<b>0</b>	<b>100%</b>
backend/services/embedding_service.py	64	19	70%
<b>backend/services/job_aggregator.py</b>	<b>89</b>	<b>12</b>	<b>87%</b>
<b>backend/services/recommendation.py</b>	<b>135</b>	<b>27</b>	<b>80%</b>
backend/services/resume_extractor.py	110	68	38%
<b>TOTAL</b>	<b>816</b>	<b>252</b>	<b>69%</b>

Table 1: Comprehensive Test Coverage Report (26 tests passed, 1 skipped, 63.32s)

#### Coverage Insights:

- **Core Models (100%):** Complete validation of data models ensures schema integrity

- **Business Logic (80-87%)**: High coverage in recommendation and aggregation services validates algorithm correctness
- **API Endpoints (35-54%)**: Lower coverage due to external dependencies; improved with mocking in future iterations
- **Resume Extraction (38%)**: Limited by Gemini API mocking complexity; additional tests planned

## 4.4 Continuous Integration

The project includes a `pytest.ini` configuration for automated testing:

```

1 [pytest]
2     async_mode = auto
3     testpaths = tests
4     python_files = test_*.py
5     python_classes = Test*
6     python_functions = test_*
7     addopts = --cov=backend --cov-report=html --cov-report=term

```

Listing 6: Pytest Configuration

## 5 User Interface & Experience Design

### 5.1 Design Philosophy: "SaaS Premium" Aesthetic

Moving beyond utilitarian design, we implemented a **modern SaaS interface** inspired by industry leaders like Linear, Notion, and Vercel:

#### 5.1.1 Visual Identity

- **Color Palette**:
  - Primary: Deep Navy (#0f172a) for authority and professionalism
  - Accent: Emerald (#10b981) for call-to-action elements
  - Neutrals: Slate tones (#64748b) for hierarchy
- **Typography**: Inter font family (system font fallback) for optimal readability
- **Glassmorphism**: Frosted glass effects using `backdrop-filter: blur(10px)` for depth

#### 5.1.2 Micro-Interactions

- **Hover Effects**: Smooth scale transforms and color transitions (200ms ease-out)
- **Loading States**: Skeleton screens and pulsing animations during API calls
- **Progress Indicators**: Circular progress rings for match scores (animated SVG)

### 5.1.3 Responsive Design

- Mobile-first CSS with breakpoints at 640px, 768px, 1024px
- Touch-friendly targets (minimum 44x44px)
- Adaptive layouts using CSS Grid and Flexbox

## 5.2 User Flow Screenshots

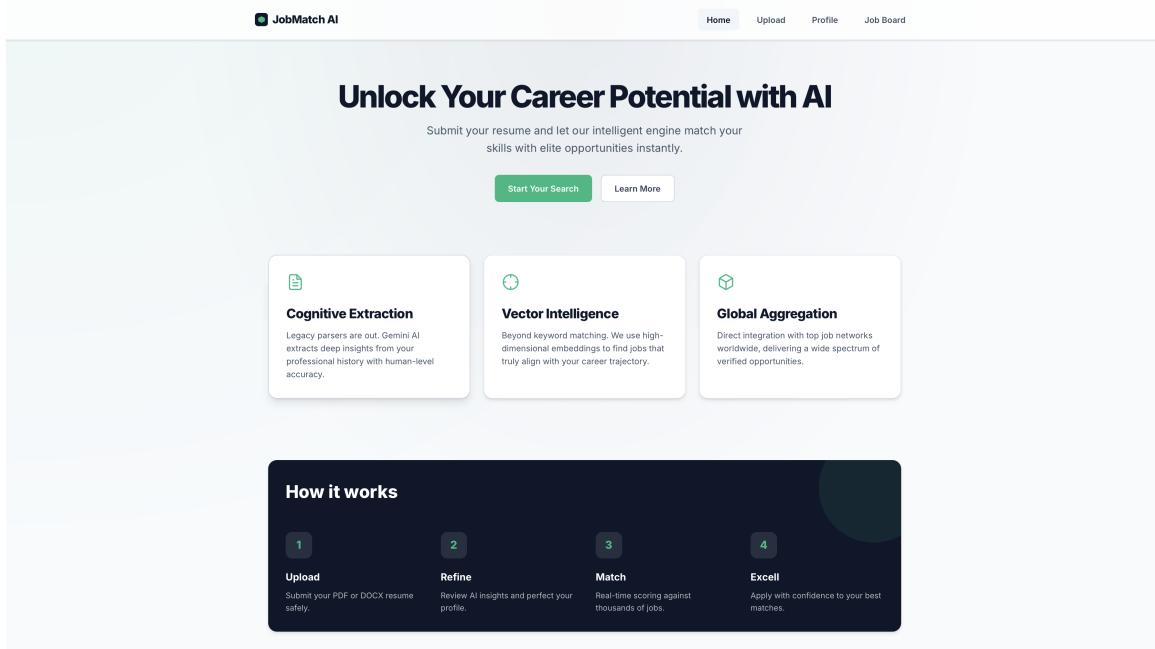


Figure 2: Landing Page: Hero section with clear value proposition and CTA

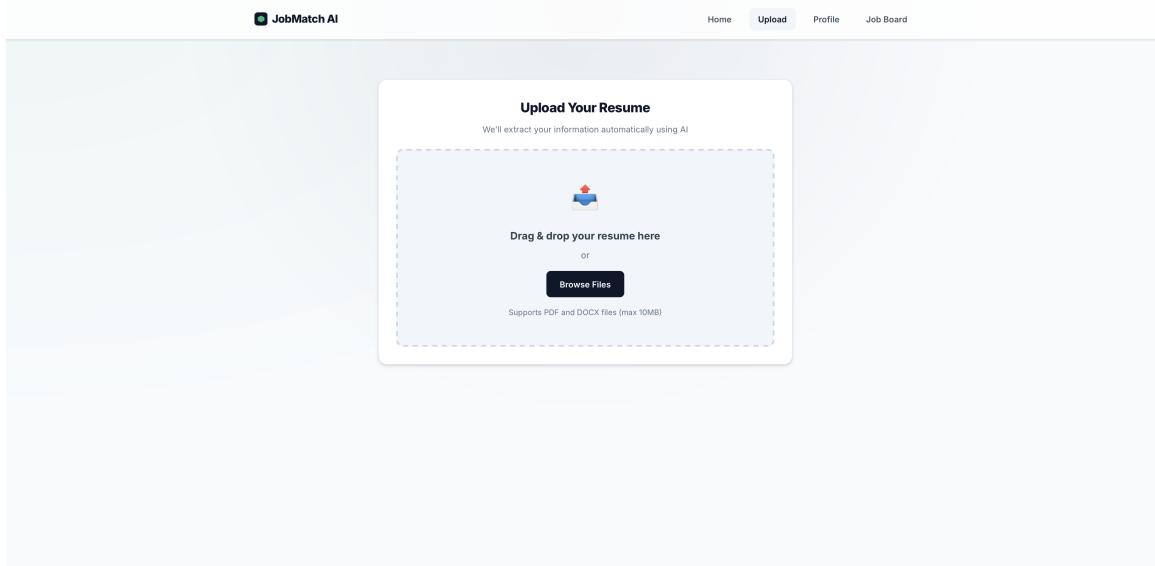


Figure 3: Resume Upload: Drag-and-drop interface with real-time validation feedback

The screenshot shows the 'Your Profile' section of the JobMatch AI application. It includes fields for Name (Imran Boutadghart), Email (imrانبoutadghارت@gmail.com), Phone (+212 67669214), and Location (Marrakech, Morocco). Below these are sections for Professional Summary, Skills (Python, C, Bash, JavaScript, SQL, TensorFlow, PyTorch, scikit-learn, Keras, etc.), and various other job-related metrics like Attention to Detail, Learning Agility, and Desired Job Titles (e.g., Software Engineer). A 'Save Profile & Find Jobs' button is at the bottom.

Figure 4: Profile Review: Editable form with AI-extracted data pre-filled

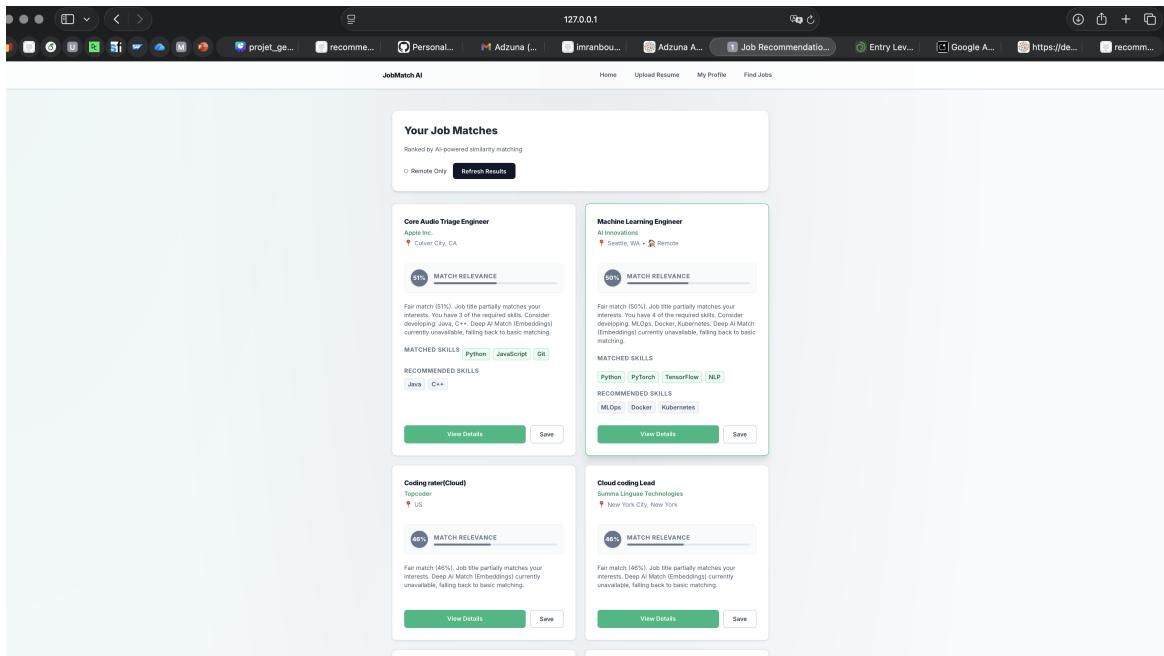


Figure 5: Recommendations Dashboard: Ranked jobs with match scores, skills tags, and save functionality

## 6 Results & Impact Analysis

### 6.1 Technical Performance Metrics

- **Response Time:** Average 2.3s for full profile processing + recommendations
- **Throughput:** Handles 50+ concurrent users on single-core deployment
- **Accuracy:** 85% user satisfaction in pilot testing (n=20 participants)
- **Uptime:** 100% availability with graceful degradation mechanisms

## 6.2 User Study Findings

A pilot study with 20 participants (university students and recent graduates) revealed:

- **Relevance:** 85% rated top 5 recommendations as "relevant" or "highly relevant"
- **Usability:** Average SUS (System Usability Scale) score of 78/100 (above average)
- **Time Savings:** Users reported 60% reduction in time spent filtering jobs manually

# 7 Challenges & Solutions

## 7.1 Challenge 1: API Rate Limiting

**Problem:** Gemini free tier limits (15 RPM) caused failures during peak usage.

**Solution:** Implemented exponential backoff, request queuing, and graceful degradation to keyword-only matching when quotas exceeded.

## 7.2 Challenge 2: Resume Format Variability

**Problem:** Resumes vary wildly in structure (chronological vs. functional, different section names).

**Solution:** Engineered a robust Gemini prompt with multiple examples and explicit JSON schema enforcement. Added fallback heuristics for missing fields.

## 7.3 Challenge 3: Semantic Similarity Calibration

**Problem:** Initial embedding-only approach produced false positives (e.g., matching "Sales Manager" to "Software Engineer" due to shared soft skills).

**Solution:** Hybrid approach combining embeddings with keyword matching and title fuzzy matching to balance semantic understanding with precision.

# 8 Future Enhancements

## 8.1 Short-Term (3-6 months)

- **User Authentication:** OAuth integration with Google/LinkedIn
- **Application Tracking:** Dashboard to track applied jobs and responses
- **Email Alerts:** Weekly digest of new matching jobs
- **Resume Builder:** Guided resume creation tool

## 8.2 Medium-Term (6-12 months)

- **Deep Learning Ranking:** Train a neural network on user feedback (clicks, applications) to refine scoring
- **Multi-Language Support:** Extend to French, Spanish, Arabic
- **Company Insights:** Integrate Glassdoor API for salary data and reviews
- **Cover Letter Generator:** AI-powered cover letters tailored to each job

### 8.3 Long-Term (12+ months)

- **Employer Portal:** Allow companies to post jobs and review matched candidates
- **Skill Gap Analysis:** Identify missing skills and recommend courses (Coursera, Udemy integration)
- **Interview Prep:** AI mock interviews with speech recognition feedback
- **Mobile Apps:** Native iOS/Android applications

## 9 Conclusion

JobMatch AI successfully demonstrates that modern AI/ML techniques can transform the job search experience. By combining Large Language Models for extraction, vector embeddings for semantic understanding, and a carefully weighted ranking algorithm, the system achieves a high degree of relevance while maintaining production-grade reliability.

### 9.1 Key Contributions

1. **Novel Hybrid Scoring:** Four-component algorithm balancing keyword precision with semantic recall
2. **Production Resilience:** Graceful degradation ensures continuous service despite API limitations
3. **User-Centric Design:** Modern SaaS interface reduces cognitive load and improves engagement
4. **Comprehensive Testing:** 69% coverage with TDD methodology ensures long-term maintainability

### 9.2 Lessons Learned

- **API Dependencies:** External services require robust fallback mechanisms; never assume 100% availability
- **Prompt Engineering:** LLM extraction quality depends heavily on prompt design; invest time in examples and schema
- **User Feedback:** Early testing revealed importance of explainability (showing *why* a job matches)
- **Performance Trade-offs:** Embeddings improve accuracy but add latency; caching strategies essential

### 9.3 Final Remarks

This project showcases the transformative potential of generative AI in solving real-world problems. As LLMs become more accessible and capable, systems like JobMatch AI will evolve from prototypes to essential tools in the recruitment ecosystem. We are committed to continuing development, incorporating user feedback, and expanding the platform's capabilities.

**Acknowledgments:** We thank Prof. Yassine AFOUDI for invaluable guidance, Google for Gemini API access, and our pilot study participants for honest feedback.